LSTM Statistics and R short course

Session 1 - Practical

Marc Henrion, Carrie Barrett

2025-03-05

Session 1 - Practical

Go to the course website on GitHub:

https://github.com/mlw-stats/R_And_Statistics_Training_LSTM2025/Session1

From here, download the following files:

```
btTBreg.csv
btTBregHospitals.csv
btTBreg_info.txt
```

1. Load the `btTBreg.csv` data table into R.

2. The variables `cd41`, `cd42` and `cd41.sk`, `cd42.sk` measure the same variables (`cd4` and `cd4.sk` respectively) in the same individuals at two different time point. This means the data are in wide format. Reformat to long format.

3. Save the reformatted data into a file called `btTBregLong.tab` in such a way that

   i. Columns are tab-separated.
   ii. Column names are saved.
   iii. No row number is saved in the resulting file.

4. Copy the code below to generate some wide-format data. We will assume this dataset contains observations of 2 biomarkers, `ferritin` and `rbp4` for 10 study participants at 2 different timepoints, `day1` and `day90`.

```
set.seed(123)

df<-data.frame(
```

```
    id=paste(sep="","P",1:10),
    ferritin_day1=rexp(10,rate=1/195),
    rbp4_day1=rexp(10,rate=1/2.5)
) %>%
    mutate(
      ferritin_day90=rnorm(10,mean=ferritin_day1+5,sd=4),
      rbp4_day90=rbp4_day1+rexp(10,rate=1/0.25)
    )
```

This is what this data table looks like:

```
   id ferritin_day1  rbp4_day1 ferritin_day90 rbp4_day90
1   P1    164.474166  2.5120751     169.251422   3.021712
2   P2    112.439003  1.2005368     114.301474   1.552174
3   P3    259.165699  0.7025341     261.231686   0.829188
4   P4      6.157585  0.9427946      10.294123   1.007684
5   P5     10.961140  0.4707101      14.621489   1.119933
6   P6     61.717737  2.1244653      62.374941   2.431722
7   P7     61.274322  3.9080088      65.932629   4.105679
8   P8     28.327027  1.1969010      37.609469   1.354221
9   P9    531.616111  1.4773371     536.034536   1.790997
10 P10      5.684922 10.1025293       6.022743  10.249700
```

Reformat this to long format, i.e. so that you have 4 columns: `id`, `time`, `ferritin` and `rbp4`.

5. Load the `btTBregHospitals.csv` data table. Join the data frames storing `btTBreg.csv` and `btTBregHospitals.csv`.

6. Compute the average patient age and the proportion of male patients for each hospital.

Useful functions for this are `aggregate()` and `group_by()`. You can however also do it manually.

7. Write an R function that computes the following summary statistics, then, using your custom function, compute these for the `bmi`, `cd41`, `cd42` columns:

    i. mean
    ii. median
    iii. interquartile range
    iv. minimum
    v. maximum
    vi. number of missing values

8. Do the same now, but only for female patients. Repeat for only male patients.